

Assignment 2 - Logistic Regression

Eoin Flynn

5 March 2018

Bond University
Data Science

Contents

Functions	3
Data	4
Split Data	5
Decision Tree	6
Logistic Regression	9
All Variables Logistic Regression	9
Top Three Logistic Regression	11
Model Comparison	13

Functions

This section will hold all of the functions that will be used throughout this markdown.

```
# Create a decision tree
createDecisionTreeModel <- function(formula, dataset, maxdepth) {
  suppressMessages(library(party))
  decisionTreeModel <- ctree(formula, data = dataset, controls = ctree_control(maxdepth = maxdepth))

  return(decisionTreeModel)
}

# Change rows to factors
setRowAsFactor <- function(dataset, columns) {
  for (column in columns) {
    dataset[, column] <- as.factor(dataset[, column])
  }
  return(dataset)
}

# Create a logistic regression model
createLogisticRegressionModel <- function(formula, family = binomial,
  dataset) {
  logisticRegressionModel = glm(formula, family = family, data = dataset)

  return(logisticRegressionModel)
}

# Create a prediction dataframe
createPrediction_df <- function(model, dataset, predictionType = "response",
  oneClass, zeroClass) {
  # Run the prediction
  prediction <- predict(model, dataset, type = predictionType)
  # Convert to a dataframe
  prediction_df <- data.frame(prediction)
  # Rename the column to reference easier
  colnames(prediction_df) <- "probabilities"
  # Add a row for the classification
  prediction_df$classification <- rep(zeroClass, nrow(prediction_df))
  # Convert all probabilities above 0.5 to be the affirmative
  # class
  prediction_df$classification[prediction_df$probabilities >
    0.5] <- oneClass
  prediction_df$classification <- as.factor(prediction_df$classification)

  return(prediction_df)
}

# Get model performance for plotting ROC curve
getModelPerformance <- function(model, dataset, outcomeColumn,
  type = "response", xAxis = "tpr", yAxis = "fpr") {
  suppressMessages(library(ROCR))
  # Create a predict variable
```

```

predict <- predict(model, dataset, type = type)
# Create a prediction variable
prediction <- prediction(predict, outcomeColumn)
# Create the performance variable
performance <- performance(prediction, xAxis, yAxis)

return(performance)
}

# Plot ROC curves
plotROCCurves <- function(model1, model2, main, model1Colour = "#009900",
  model2Colour = "#FF8000", model1Name, model2Name, legendLocation = "bottomright") {
  plot(model1, main = main, col = model1Colour, print.auc = TRUE)
  plot(model2, add = T, col = model2Colour)
  legend(legendLocation, legend = paste(rep(c(model1Name, model2Name))),
    col = c(model1Colour, model2Colour), cex = 0.8, fill = c(model1Colour,
    model2Colour))
}

# Calculate AUC. Returns as a decimal
getAUC <- function(outcomeColumn, dataset, model, oneClass, zeroClass) {
  suppressMessages(library(ModelMetrics))
  prediction_df <- createPrediction_df(model, dataset, oneClass = oneClass,
    zeroClass = zeroClass)
  auc <- auc(outcomeColumn, prediction_df$classification)

  return(auc)
}

# Create a confusion matrix
createConfusionMatrix <- function(model, dataset, outcomeColumn,
  oneClass, zeroClass) {
  suppressMessages(library(caret))
  prediction_df <- createPrediction_df(model, dataset, oneClass = oneClass,
    zeroClass = zeroClass)
  userConfusionMatrix <- table(outcomeColumn, prediction_df$classification)

  return(userConfusionMatrix)
}

```

Data

In this section we will load in our data and do some basic data exploration.

```

suppressMessages(library(RMySQL))

USER <- "root"
PASSWORD <- "A13337995"
HOST <- "localhost"
DBNAME <- "world"

statement <- "Select * from world.customerChurn"

```

```

db <- dbConnect(MySQL(), user = USER, password = PASSWORD, host = HOST,
  dbname = DBNAME, port = 3306)
customerDataset <- dbGetQuery(db, statement = statement)
dbDisconnect(db)

## [1] TRUE

# Loops through and changes all relevant rows to factors and
# returns the dataset post modification

customerDataset <- setRowAsFactor(customerDataset, c("gender",
  "SeniorCitizen", "Partner", "Dependents", "PhoneService",
  "MultipleLines", "InternetService", "OnlineSecurity", "OnlineBackup",
  "DeviceProtection", "TechSupport", "StreamingTV", "StreamingMovies",
  "Contract", "PaperlessBilling", "PaymentMethod", "Churn"))

# Drop the columns that will not be needed
customerDataset = customerDataset[, -which(names(customerDataset) %in%
  c("customerID"))]

```

Split Data

We will now split our data into test and training sets. The purpose of this is to create a sample of data that the model has never seen before in order to gauge its accuracy. The training set will consist of 80% of the data while the remaining 20% will constitute the test set.

```

suppressMessages(library(caTools))

# Set the seed to reproducibility
set.seed(12216)

# Create our two datasets
sample <- sample.split(customerDataset, SplitRatio = 0.8)
train_df <- subset(customerDataset, sample == TRUE)
test_df <- subset(customerDataset, sample == FALSE)

# We can now see that the data is split approximately 80:20
print(sprintf("The full dataset has %s observations", NROW(customerDataset)))

## [1] "The full dataset has 7032 observations"

print(sprintf("The training dataset has %s observations", NROW(train_df)))

## [1] "The training dataset has 5628 observations"

print(sprintf("The testing dataset has %s observations", NROW(test_df)))

## [1] "The testing dataset has 1404 observations"

# Check to see how many customers churned in each dataset
table(train_df$Churn)

##
##    No    Yes
## 4120 1508

```

```

table(test_df$Churn)

##
##   No   Yes
## 1043  361
# We can see that each dataset holds approximately the same
# proportion of customers who churned
print(sprintf("%.2f%% of the training set churned", ((NROW(subset(train_df,
  Churn == "Yes")))/NROW(train_df) * 100)))

## [1] "26.79% of the training set churned"
print(sprintf("%.2f%% of the testing set churned", ((NROW(subset(test_df,
  Churn == "Yes")))/NROW(test_df) * 100)))

## [1] "25.71% of the testing set churned"

```

Decision Tree

We want to first make a decision tree to determine which variables are best able to predict whether a customer will churn. We already know from our previous report that the optimal model is however this time the tree will only be run on the training dataset.

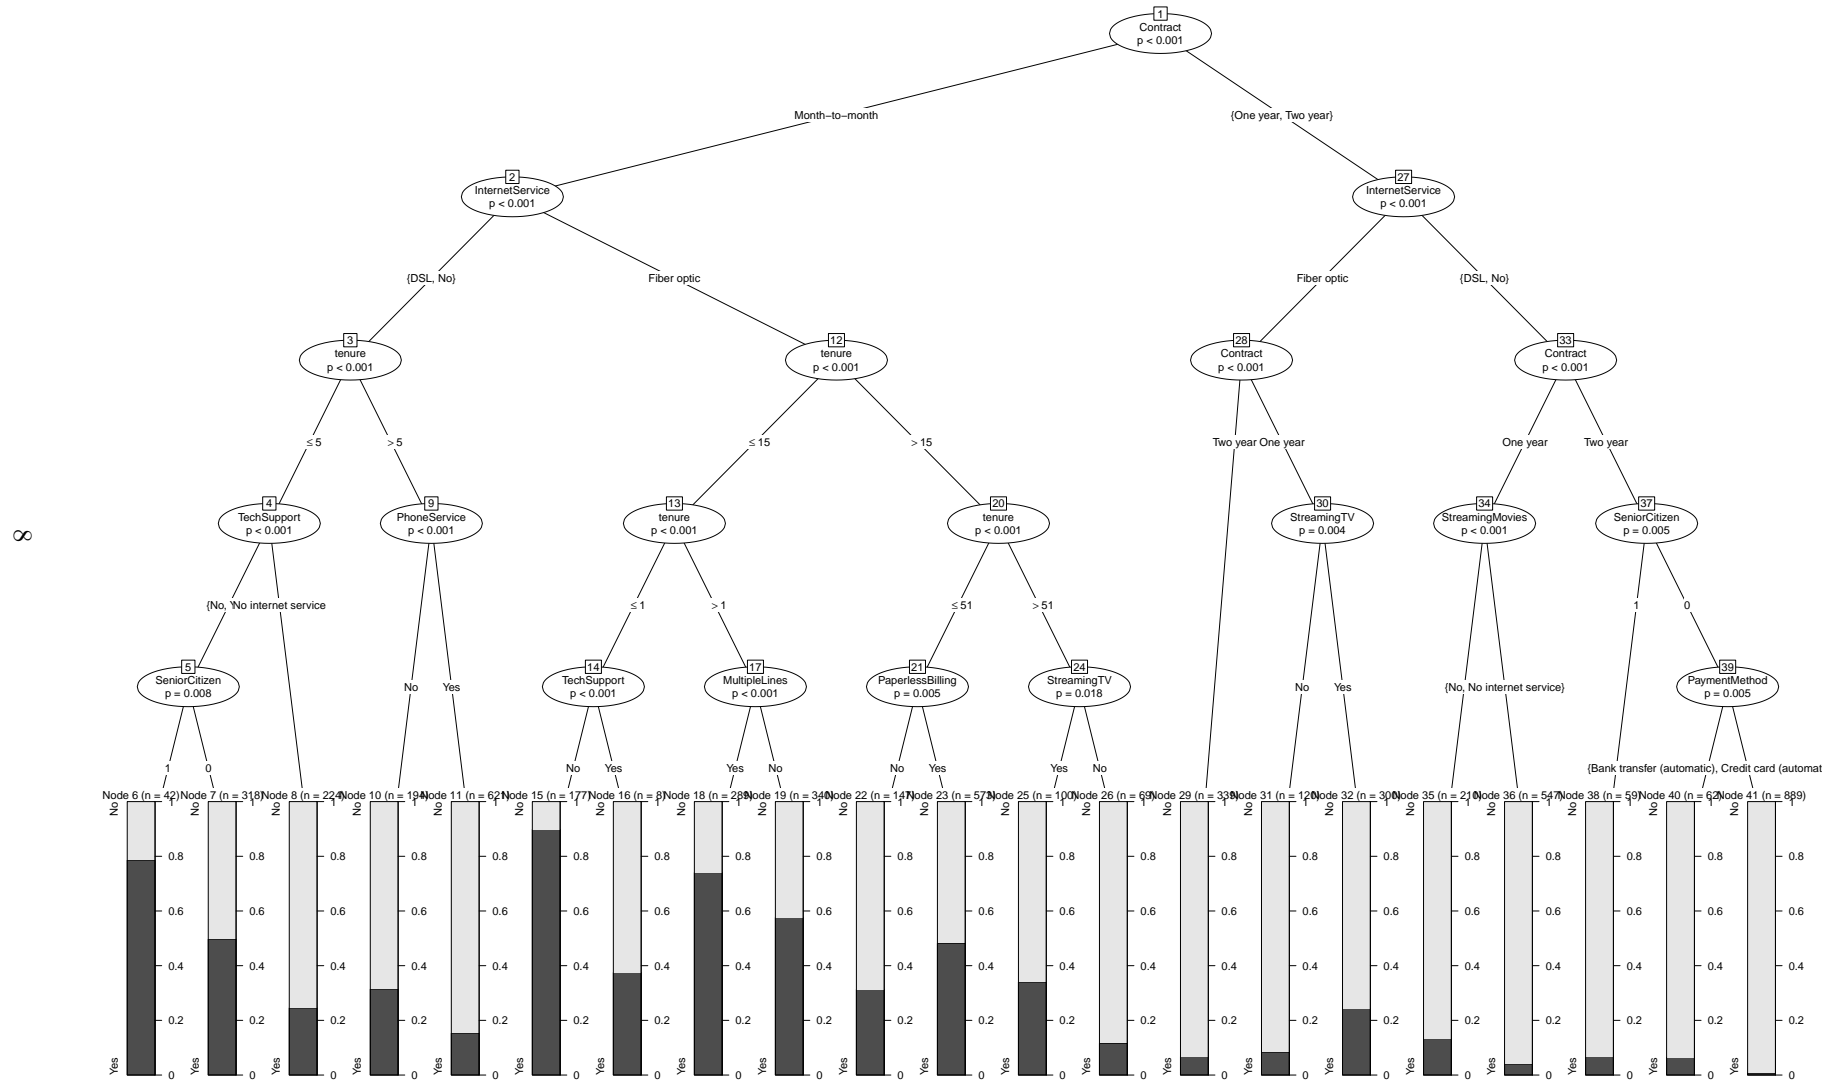
```

# Create and plot the decision tree
decisionTreeModel = createDecisionTreeModel(formula = Churn ~
  ., dataset = train_df, maxdepth = 5)

```

```
plot(decisionTreeModel, main = "Decision Tree Model", type = "extended",  
     newpage = TRUE)
```

Decision Tree Model



From looking at the decision tree it is clear that the top three variables are Contract, InternetService, and Tenure. Below those three levels, other variables such as StreamingTV, and TechSupport become relevant. To develop the best performing model with this dataset we will create a regression using only those three top level variables, and then a second using all variables. The results from each model will then be compared before the best model is presented to management.

Logistic Regression

We will now create multiple logistic regressions using GLM package. We will compare and then optimize before providing a final model to management for business use.

All Variables Logistic Regression

In this section we will create a logistic regression using all variables in the training dataset and look at statistical significance of each. Later in the report we will assess the accuracy of this model against others.

```
# We will start by first making a regression using all
# variables
allVariablesLogisticRegressionModel <- createLogisticRegressionModel(formula = Churn ~
  ., dataset = train_df)

# Print a summary of the regression
print("Model Summary")
```

```
## [1] "Model Summary"
```

```
summary(allVariablesLogisticRegressionModel)
```

```
##
## Call:
## glm(formula = formula, family = family, data = dataset)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9271  -0.6763  -0.2846   0.7434   3.4193
##
## Coefficients: (7 not defined because of singularities)
##
##              Estimate Std. Error z value
## (Intercept)    1.168e+00  9.050e-01  1.290
## genderMale      5.431e-02  7.237e-02  0.750
## SeniorCitizen1  2.279e-01  9.429e-02  2.417
## PartnerYes      3.171e-02  8.752e-02  0.362
## DependentsYes  -1.351e-01  1.002e-01 -1.348
## tenure         -5.790e-02  6.865e-03 -8.434
## PhoneServiceYes 1.985e-01  7.223e-01  0.275
## MultipleLinesNo phone service      NA         NA      NA
## MultipleLinesYes  4.521e-01  1.974e-01  2.290
## InternetServiceFiber optic    1.765e+00  8.883e-01  1.987
## InternetServiceNo -1.900e+00  8.981e-01 -2.116
## OnlineSecurityNo internet service      NA         NA      NA
## OnlineSecurityYes -2.001e-01  1.994e-01 -1.003
## OnlineBackupNo internet service      NA         NA      NA
## OnlineBackupYes   6.721e-02  1.948e-01  0.345
```

```

## DeviceProtectionNo internet service      NA      NA      NA
## DeviceProtectionYes                      1.476e-01 1.965e-01 0.751
## TechSupportNo internet service           NA      NA      NA
## TechSupportYes                          -1.518e-01 2.011e-01 -0.755
## StreamingTVNo internet service           NA      NA      NA
## StreamingTVYes                          6.178e-01 3.625e-01 1.704
## StreamingMoviesNo internet service       NA      NA      NA
## StreamingMoviesYes                      5.822e-01 3.650e-01 1.595
## ContractOne year                        -6.773e-01 1.197e-01 -5.657
## ContractTwo year                       -1.424e+00 2.001e-01 -7.114
## PaperlessBillingYes                    3.382e-01 8.284e-02 4.083
## PaymentMethodCredit card (automatic) -5.187e-02 1.273e-01 -0.407
## PaymentMethodElectronic check          3.014e-01 1.056e-01 2.855
## PaymentMethodMailed check              -4.768e-02 1.280e-01 -0.372
## MonthlyCharges                        -4.160e-02 3.531e-02 -1.178
## TotalCharges                          2.991e-04 7.846e-05 3.813
## Pr(>|z|)
## (Intercept)                          0.197001
## genderMale                          0.452970
## SeniorCitizen1                      0.015645 *
## PartnerYes                          0.717142
## DependentsYes                      0.177760
## tenure                             < 2e-16 ***
## PhoneServiceYes                    0.783462
## MultipleLinesNo phone service         NA
## MultipleLinesYes                   0.022016 *
## InternetServiceFiber optic          0.046872 *
## InternetServiceNo                   0.034383 *
## OnlineSecurityNo internet service     NA
## OnlineSecurityYes                   0.315685
## OnlineBackupNo internet service       NA
## OnlineBackupYes                     0.730109
## DeviceProtectionNo internet service     NA
## DeviceProtectionYes                   0.452539
## TechSupportNo internet service         NA
## TechSupportYes                       0.450396
## StreamingTVNo internet service         NA
## StreamingTVYes                       0.088343 .
## StreamingMoviesNo internet service     NA
## StreamingMoviesYes                   0.110691
## ContractOne year                     1.54e-08 ***
## ContractTwo year                     1.13e-12 ***
## PaperlessBillingYes                  4.45e-05 ***
## PaymentMethodCredit card (automatic) 0.683705
## PaymentMethodElectronic check         0.004305 **
## PaymentMethodMailed check            0.709541
## MonthlyCharges                       0.238771
## TotalCharges                         0.000138 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6542.0 on 5627 degrees of freedom

```

```
## Residual deviance: 4674.1 on 5604 degrees of freedom
## AIC: 4722.1
##
## Number of Fisher Scoring iterations: 6
# Convert Betas into odds ratio
print("Odds Ratios")

## [1] "Odds Ratios"
exp(coef(allVariablesLogisticRegressionModel))

##              (Intercept)              genderMale
##              3.2142616              1.0558152
##              SeniorCitizen1              PartnerYes
##              1.2559792              1.0322134
##              DependentsYes              tenure
##              0.8736378              0.9437414
##              PhoneServiceYes      MultipleLinesNo phone service
##              1.2195730              NA
##              MultipleLinesYes      InternetServiceFiber optic
##              1.5716412              5.8440322
##              InternetServiceNo      OnlineSecurityNo internet service
##              0.1495586              NA
##              OnlineSecurityYes      OnlineBackupNo internet service
##              0.8186677              NA
##              OnlineBackupYes      DeviceProtectionNo internet service
##              1.0695253              NA
##              DeviceProtectionYes      TechSupportNo internet service
##              1.1590444              NA
##              TechSupportYes      StreamingTVNo internet service
##              0.8591911              NA
##              StreamingTVYes      StreamingMoviesNo internet service
##              1.8547853              NA
##              StreamingMoviesYes      ContractOne year
##              1.7900570              0.5079813
##              ContractTwo year      PaperlessBillingYes
##              0.2408675              1.4024619
##      PaymentMethodCredit card (automatic)      PaymentMethodElectronic check
##              0.9494527              1.3517921
##              PaymentMethodMailed check      MonthlyCharges
##              0.9534397              0.9592580
##              TotalCharges
##              1.0002992
```

From the model's summary we can see that our top three variables identified earlier are all have a high statistical significance. TotalCharges is also statistically different from zero however since the decision tree deemed that it had no predictive information we will not be including it in our model building

Top Three Logistic Regression

We will now create our logistic regression using only the top three variables from our decision tree (Contract, InternetService, and Tenure)

```

# We will start by first making a regression using all
# variables
topThreeLogisticRegressionModel <- createLogisticRegressionModel(formula = Churn ~
  Contract + InternetService + tenure, dataset = train_df)

# Print a summary of the regression
print("Model Summary")

## [1] "Model Summary"

summary(topThreeLogisticRegressionModel)

##
## Call:
## glm(formula = formula, family = family, data = dataset)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5640  -0.6746  -0.3077   0.8350   3.1500
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -0.257388    0.070388  -3.657 0.000255 ***
## ContractOne year    -0.870290    0.113829  -7.646 2.08e-14 ***
## ContractTwo year   -1.740473    0.191348  -9.096 < 2e-16 ***
## InternetServiceFiber optic  1.163022    0.080631  14.424 < 2e-16 ***
## InternetServiceNo   -1.121017    0.131645  -8.515 < 2e-16 ***
## tenure            -0.031107    0.002167 -14.353 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6542.0  on 5627  degrees of freedom
## Residual deviance: 4852.3  on 5622  degrees of freedom
## AIC: 4864.3
##
## Number of Fisher Scoring iterations: 6

# Convert Betas into odds ratio
print("Odds Ratios")

## [1] "Odds Ratios"

exp(coef(topThreeLogisticRegressionModel))

##              (Intercept)          ContractOne year
##              0.7730685              0.4188300
##      ContractTwo year InternetServiceFiber optic
##              0.1754374              3.1995868
##      InternetServiceNo              tenure
##              0.3259482              0.9693719

```

Model Comparison


Now that we have created all three models we will test their accuracy against the training dataset and also the test dataset. One of the easiest ways to see which model is more accurate is to use an ROC curve and measure the area under each curve, the larger the area under the curve the more accurate the model is. We will first test each model using the training dataset before testing each one individually with the test dataset to gauge how robust they are. Once we have established whether the models are robust we will create a pair of confusion matrices using the test dataset to see which model had the highest percentage of churns predicted correctly.

```
suppressMessages(library(caret))
suppressMessages(library(ROCR))

# Get the performance of each model
allVariableModelPerformance <- getModelPerformance(allVariablesLogisticRegressionModel,
  dataset = train_df, outcomeColumn = train_df$Churn)

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
topThreeModelPerformance <- getModelPerformance(topThreeLogisticRegressionModel,
  dataset = train_df, outcomeColumn = train_df$Churn)

# Plot the ROC curves
plotROCCurves(allVariableModelPerformance, topThreeModelPerformance,
  main = "ROC Curves Comparison", model1Name = "AllVariableModel",
  model2Name = "TopThreeModel")
```



```

# Get the AUC
allVariableModelAUC <- getAUC(train_df$Churn, train_df, allVariablesLogisticRegressionModel,
  oneClass = "Yes", zeroClass = "No")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
topThreeModelAUC <- getAUC(train_df$Churn, train_df, topThreeLogisticRegressionModel,
  oneClass = "Yes", zeroClass = "No")
# Print the AUC
print(sprintf("The AUC for the allVariableModel is %.2f%%", allVariableModelAUC *
  100))


## [1] "The AUC for the allVariableModel is 73.03%"
print(sprintf("The AUC for the topThreeModel is %.2f%%", topThreeModelAUC *
  100))

## [1] "The AUC for the topThreeModel is 69.21%"
# Check the allVariableModel for overfitting
allVariableModelPerformanceTest <- getModelPerformance(allVariablesLogisticRegressionModel,
  dataset = test_df, outcomeColumn = test_df$Churn)

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
plotROCCurves(allVariableModelPerformance, allVariableModelPerformanceTest,
  main = "All Variable Robustness", model1Name = "Training Dataset",
  model2Name = "Test Dataset")

```

All Variable Robustness



```
allVariableModelAUCTest <- getAUC(test_df$Churn, test_df, allVariablesLogisticRegressionModel,  
  oneClass = "Yes", zeroClass = "No")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =  
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
print(sprintf("The AUC for the training dataset using the allVariableModel is %.2f%%",  
  allVariableModelAUC * 100))
```

```
## [1] "The AUC for the training dataset using the allVariableModel is 73.03%"
```

```
print(sprintf("The AUC for the test dataset using the allVariableModel is %.2f%%",  
  allVariableModelAUCTest * 100))
```


```
## [1] "The AUC for the test dataset using the allVariableModel is 72.13%"
```

```
# Check the topThreeModel for overfitting
```

```
topThreeModelPerformanceTest <- getModelPerformance(topThreeLogisticRegressionModel,  
  dataset = test_df, outcomeColumn = test_df$Churn)
```

```
plotROCCurves(topThreeModelPerformance, topThreeModelPerformanceTest,  
  main = "Top Three Variable Robustness", model1Name = "Training Dataset",  
  model2Name = "Test Dataset")
```

Top Three Variable Robustness



```
topThreeModelAUCTest <- getAUC(test_df$Churn, test_df, topThreeLogisticRegressionModel,
  oneClass = "Yes", zeroClass = "No")
print(sprintf("The AUC for the training dataset using the topThreeModel is %.2f%%",
  topThreeModelAUC * 100))
```

```
## [1] "The AUC for the training dataset using the topThreeModel is 69.21%"
```

```
print(sprintf("The AUC for the test dataset using the topThreeModel is %.2f%%",
  topThreeModelAUCTest * 100))
```

```
## [1] "The AUC for the test dataset using the topThreeModel is 69.97%"
```

```
# Create confusion matrices to compare predictions
print("All Variable Model Confusion Matrix")
```

```
## [1] "All Variable Model Confusion Matrix"
```

```
createConfusionMatrix(allVariablesLogisticRegressionModel, test_df,
  test_df$Churn, oneClass = "Yes", zeroClass = "No")
```

```
## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading
```

```
##
## outcomeColumn No Yes
##               No  924 119
##               Yes  160 201
```

```
print("Top Three Variable Confusion Matrix")
```



```
## [1] "Top Three Variable Confusion Matrix"
print(createConfusionMatrix(topThreeLogisticRegressionModel,
  test_df, test_df$Churn, oneClass = "Yes", zeroClass = "No"))

##
## outcomeColumn  No Yes
##               No  928 115
##               Yes  177 184
```