

Assignment 2 - Logistic Regression

Eoin Flynn

5 March 2018

Bond University
Data Science

Contents

Functions	3
Data	3
Split Data	4
Decision Tree	5
Logistic Regression	8
All Variables Logistic Regression	8

Functions

This section will hold all of the functions that will be used throughout this markdown.

```
# Create a decision tree
createDecisionTreeModel <- function(formula, dataset, maxdepth) {
  suppressMessages(library(party))
  decisionTreeModel <- ctree(formula, data = dataset, controls = ctree_control(maxdepth = maxdepth))

  return(decisionTreeModel)
}

# Change rows to factors
setRowAsFactor <- function(dataset, columns) {
  for (column in columns) {
    dataset[, column] <- as.factor(dataset[, column])
  }
  return(dataset)
}

# Create a logistic regression model
createLogisticRegressionModel <- function(formula, family = binomial,
  dataset) {
  logisticRegressionModel = glm(formula, family = family, data = dataset)

  return(logisticRegressionModel)
}
```

Data

In this section we will load in our data and do some basic data exploration.

```
suppressMessages(library(RMySQL))

USER <- "root"
PASSWORD <- "A13337995"
HOST <- "localhost"
DBNAME <- "world"

statement <- "Select * from world.customerChurn"
db <- dbConnect(MySQL(), user = USER, password = PASSWORD, host = HOST,
  dbname = DBNAME, port = 3306)
customerDataset <- dbGetQuery(db, statement = statement)
dbDisconnect(db)
```

```
## [1] TRUE
```

```
# Loops through and changes all relevant rows to factors and
# returns the dataset post modification
```

```
customerDataset <- setRowAsFactor(customerDataset, c("gender",
  "SeniorCitizen", "Partner", "Dependents", "PhoneService",
  "MultipleLines", "InternetService", "OnlineSecurity", "OnlineBackup",
```

```

    "DeviceProtection", "TechSupport", "StreamingTV", "StreamingMovies",
    "Contract", "PaperlessBilling", "PaymentMethod", "Churn"))

# Drop the columns that will not be needed
customerDataset = customerDataset[, -which(names(customerDataset) %in%
    c("customerID"))]

```

Split Data

We will now split our data into test and training sets. The purpose of this is to create a sample of data that the model has never seen before in order to gauge its accuracy. The training set will consist of 80% of the data while the remaining 20% will constitute the test set.

```

suppressMessages(library(caTools))

# Set the seed to reproducibility
set.seed(12216)

# Create our two datasets
sample <- sample.split(customerDataset, SplitRatio = 0.8)
train_df <- subset(customerDataset, sample == TRUE)
test_df <- subset(customerDataset, sample == FALSE)

# We can now see that the data is split approximately 80:20
print(sprintf("The full dataset has %s observations", NROW(customerDataset)))

## [1] "The full dataset has 7032 observations"

print(sprintf("The training dataset has %s observations", NROW(train_df)))

## [1] "The training dataset has 5628 observations"

print(sprintf("The testing dataset has %s observations", NROW(test_df)))

## [1] "The testing dataset has 1404 observations"

# Check to see how many customers churned in each dataset
table(train_df$Churn)

##
##   No  Yes
## 4120 1508

table(test_df$Churn)

##
##   No  Yes
## 1043  361

# We can see that each dataset holds approximately the same
# proportion of customers who churned
print(sprintf("%.2f%% of the training set churned", ((NROW(subset(train_df,
    Churn == "Yes")))/NROW(train_df) * 100)))

## [1] "26.79% of the training set churned"

print(sprintf("%.2f%% of the testing set churned", ((NROW(subset(test_df,
    Churn == "Yes")))/NROW(test_df) * 100)))

```

```
## [1] "25.71% of the testing set churned"
```


Decision Tree

We want to first make a decision tree to determine which variables are best able to predict whether a customer will churn. We already know from our previous report that the optimal model is however this time the tree will only be run on the training dataset.

```
# Create and plot the decision tree  
decisionTreeModel = createDecisionTreeModel(formula = Churn ~  
  ., dataset = train_df, maxdepth = 5)
```

```
plot(decisionTreeModel, main = "Decision Tree Model", type = "extended",  
     newpage = TRUE)
```

Decision Tree Model



From looking at the decision tree it is clear that the top three variables are Contract, InternetService, and Tenure. Below those three levels, other variables such as StreamingTV, and TechSupport become relevant. To develop the best performing model with this dataset we will create a regression using only those three top level variables, and then a second using all variables. The results from each model will then be compared before the best model is presented to management.

Logistic Regression

We will now create multiple logistic regressions using GLM package. We will compare and then optimize before providing a final model to management for business use.

All Variables Logistic Regression

In this section we will create a logistic regression using all variables in the training dataset and look at statistical significance of each. Later in the report we will assess the accuracy of this model against others.

```
# We will start by first making a regression using all
# variables
allVariablesLogisticRegressionModel <- createLogisticRegressionModel(formula = Churn ~
  ., dataset = train_df)

# Print a summary of the regression
print("Model Summary")
```

```
## [1] "Model Summary"
```

```
summary(allVariablesLogisticRegressionModel)
```

```
##
## Call:
## glm(formula = formula, family = family, data = dataset)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.9271  -0.6763  -0.2846   0.7434   3.4193
##
## Coefficients: (7 not defined because of singularities)
##
##              Estimate Std. Error z value
## (Intercept)    1.168e+00  9.050e-01  1.290
## genderMale      5.431e-02  7.237e-02  0.750
## SeniorCitizen1  2.279e-01  9.429e-02  2.417
## PartnerYes      3.171e-02  8.752e-02  0.362
## DependentsYes  -1.351e-01  1.002e-01 -1.348
## tenure         -5.790e-02  6.865e-03 -8.434
## PhoneServiceYes 1.985e-01  7.223e-01  0.275
## MultipleLinesNo phone service      NA         NA      NA
## MultipleLinesYes  4.521e-01  1.974e-01  2.290
## InternetServiceFiber optic    1.765e+00  8.883e-01  1.987
## InternetServiceNo -1.900e+00  8.981e-01 -2.116
## OnlineSecurityNo internet service      NA         NA      NA
## OnlineSecurityYes -2.001e-01  1.994e-01 -1.003
## OnlineBackupNo internet service      NA         NA      NA
## OnlineBackupYes   6.721e-02  1.948e-01  0.345
```



```

## DeviceProtectionNo internet service      NA      NA      NA
## DeviceProtectionYes                      1.476e-01 1.965e-01 0.751
## TechSupportNo internet service           NA      NA      NA
## TechSupportYes                          -1.518e-01 2.011e-01 -0.755
## StreamingTVNo internet service           NA      NA      NA
## StreamingTVYes                          6.178e-01 3.625e-01 1.704
## StreamingMoviesNo internet service       NA      NA      NA
## StreamingMoviesYes                      5.822e-01 3.650e-01 1.595
## ContractOne year                        -6.773e-01 1.197e-01 -5.657
## ContractTwo year                       -1.424e+00 2.001e-01 -7.114
## PaperlessBillingYes                    3.382e-01 8.284e-02 4.083
## PaymentMethodCredit card (automatic) -5.187e-02 1.273e-01 -0.407
## PaymentMethodElectronic check          3.014e-01 1.056e-01 2.855
## PaymentMethodMailed check              -4.768e-02 1.280e-01 -0.372
## MonthlyCharges                        -4.160e-02 3.531e-02 -1.178
## TotalCharges                          2.991e-04 7.846e-05 3.813
## Pr(>|z|)
## (Intercept)                          0.197001
## genderMale                          0.452970
## SeniorCitizen1                      0.015645 *
## PartnerYes                          0.717142
## DependentsYes                      0.177760
## tenure                             < 2e-16 ***
## PhoneServiceYes                    0.783462
## MultipleLinesNo phone service         NA
## MultipleLinesYes                   0.022016 *
## InternetServiceFiber optic          0.046872 *
## InternetServiceNo                  0.034383 *
## OnlineSecurityNo internet service     NA
## OnlineSecurityYes                  0.315685
## OnlineBackupNo internet service       NA
## OnlineBackupYes                   0.730109
## DeviceProtectionNo internet service     NA
## DeviceProtectionYes                 0.452539
## TechSupportNo internet service         NA
## TechSupportYes                     0.450396
## StreamingTVNo internet service         NA
## StreamingTVYes                     0.088343 .
## StreamingMoviesNo internet service     NA
## StreamingMoviesYes                  0.110691
## ContractOne year                   1.54e-08 ***
## ContractTwo year                   1.13e-12 ***
## PaperlessBillingYes                4.45e-05 ***
## PaymentMethodCredit card (automatic) 0.683705
## PaymentMethodElectronic check        0.004305 **
## PaymentMethodMailed check           0.709541
## MonthlyCharges                    0.238771
## TotalCharges                      0.000138 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 6542.0 on 5627 degrees of freedom

```

```
## Residual deviance: 4674.1 on 5604 degrees of freedom
## AIC: 4722.1
##
## Number of Fisher Scoring iterations: 6
```

```
# Convert Betas into odds ratio
print("Odds Ratios")
```

```
## [1] "Odds Ratios"
```

```
exp(coef(allVariablesLogisticRegressionModel))
```

```
##              (Intercept)              genderMale
##              3.2142616              1.0558152
##              SeniorCitizen1              PartnerYes
##              1.2559792              1.0322134
##              DependentsYes              tenure
##              0.8736378              0.9437414
##              PhoneServiceYes      MultipleLinesNo phone service
##              1.2195730              NA
##              MultipleLinesYes      InternetServiceFiber optic
##              1.5716412              5.8440322
##              InternetServiceNo      OnlineSecurityNo internet service
##              0.1495586              NA
##              OnlineSecurityYes      OnlineBackupNo internet service
##              0.8186677              NA
##              OnlineBackupYes      DeviceProtectionNo internet service
##              1.0695253              NA
##              DeviceProtectionYes      TechSupportNo internet service
##              1.1590444              NA
##              TechSupportYes      StreamingTVNo internet service
##              0.8591911              NA
##              StreamingTVYes      StreamingMoviesNo internet service
##              1.8547853              NA
##              StreamingMoviesYes      ContractOne year
##              1.7900570              0.5079813
##              ContractTwo year      PaperlessBillingYes
##              0.2408675              1.4024619
## PaymentMethodCredit card (automatic)      PaymentMethodElectronic check
##              0.9494527              1.3517921
##              PaymentMethodMailed check      MonthlyCharges
##              0.9534397              0.9592580
##              TotalCharges
##              1.0002992
```

From the model's summary we can see that our top three variables identified earlier are all have a high statistical significance. The surprising variable here is TotalCharges which also has a high statistical significance but was not featured on the decision tree. To evaluate what information can be gained from TotalCharges a third model will be created in which it will feature.